



## Module 9.4 : Better initialization strategies

## Deep Learning has evolved

- Better optimization algorithms
- Better regularization methods
- Better activation functions
- **Better weight initialization strategies**




- What happens if we initialize all weights to 0?



$$a_{11} = w_{11}x_1 + w_{12}x_2$$


- What happens if we initialize all weights to 0?



- What happens if we initialize all weights to 0?

$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$




- What happens if we initialize all weights to 0?

$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$




- What happens if we initialize all weights to 0?
- All neurons in layer 1 will get the same activation

$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$



- Now what will happen during back propagation?


$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$





- Now what will happen during back propagation?


$$\nabla w_{11} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{11}} \cdot \frac{\partial h_{11}}{\partial a_{11}} \cdot x_1$$

$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$



- Now what will happen during back propagation?

$$\nabla w_{11} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{11}} \cdot \frac{\partial h_{11}}{\partial a_{11}} \cdot x_1$$


$$\nabla w_{21} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{12}} \cdot \frac{\partial h_{12}}{\partial a_{12}} \cdot x_1$$

$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$



$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$


$$\therefore h_{11} = h_{12}$$

- Now what will happen during back propagation?

$$\nabla w_{11} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{11}} \cdot \frac{\partial h_{11}}{\partial a_{11}} \cdot x_1$$

$$\nabla w_{21} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{12}} \cdot \frac{\partial h_{12}}{\partial a_{12}} \cdot x_1$$

$$\text{but } h_{11} = h_{12}$$



$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$


- Now what will happen during back propagation?

$$\nabla w_{11} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{11}} \cdot \frac{\partial h_{11}}{\partial a_{11}} \cdot x_1$$

$$\nabla w_{21} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{12}} \cdot \frac{\partial h_{12}}{\partial a_{12}} \cdot x_1$$

$$\text{but } h_{11} = h_{12}$$

$$\text{and } a_{12} = a_{12}$$



$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$

- Now what will happen during back propagation?


$$\nabla w_{11} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{11}} \cdot \frac{\partial h_{11}}{\partial a_{11}} \cdot x_1$$

$$\nabla w_{21} = \frac{\partial \mathcal{L}(\mathbf{w})}{\partial y} \cdot \frac{\partial y}{\partial h_{12}} \cdot \frac{\partial h_{12}}{\partial a_{12}} \cdot x_1$$

$$\text{but } h_{11} = h_{12}$$

$$\text{and } a_{12} = a_{12}$$

$$\therefore \nabla w_{11} = \nabla w_{21}$$




- Hence both the weights will get the same update and remain equal

$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$




- Hence both the weights will get the same update and remain equal
- Infact this symmetry will never break during training

$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$



- Hence both the weights will get the same update and remain equal
- Infact this symmetry will never break during training
- The same is true for  $w_{12}$  and  $w_{22}$


$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$






$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$

- Hence both the weights will get the same update and remain equal
- Infact this symmetry will never break during training
- The same is true for  $w_{12}$  and  $w_{22}$
- And for all weights in layer 2 (infact, work out the math and convince yourself that all the weights in this layer will remain equal )




$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$

- Hence both the weights will get the same update and remain equal
- Infact this symmetry will never break during training
- The same is true for  $w_{12}$  and  $w_{22}$
- And for all weights in layer 2 (infact, work out the math and convince yourself that all the weights in this layer will remain equal )
- This is known as the **symmetry breaking problem**



$$a_{11} = w_{11}x_1 + w_{12}x_2$$

$$a_{12} = w_{21}x_1 + w_{22}x_2$$

$$\therefore a_{11} = a_{12} = 0$$

$$\therefore h_{11} = h_{12}$$

- Hence both the weights will get the same update and remain equal
- Infact this symmetry will never break during training
- The same is true for  $w_{12}$  and  $w_{22}$
- And for all weights in layer 2 (infact, work out the math and convince yourself that all the weights in this layer will remain equal )
- This is known as the **symmetry breaking problem**
- This will happen if all the weights in a network are initialized to the **same value**

We will now consider a feedforward network with:

```
D = np.random.randn(1000,500)
hidden_layer_sizes = [500]*10
nonlinearities = ['tanh']*len(hidden_layer_sizes)

act = {'relu':lambda x: np.maximum(0, x), 'tanh': lambda x: np.tanh(x),
       'sigmoid':lambda x: 1/(1 + np.exp(-x))}
Hs = {}

for i in xrange(len(hidden_layer_sizes)):
    X = D if i == 0 else Hs[i-1]
    fan_in = X.shape[1]
    fan_out = hidden_layer_sizes[i]
    W = np.random.randn(fan_in, fan_out) * 0.01

    H = np.dot(X, W)
    H = act[nonlinearities[i]](H)
    Hs[i] = H
```

```

D = np.random.randn(1000,500)
hidden_layer_sizes = [500]*10
nonlinearities = ['tanh']*len(hidden_layer_sizes)

act = {'relu':lambda x: np.maximum(0, x), 'tanh': lambda x: np.tanh(x),
       'sigmoid':lambda x: 1/(1 + np.exp(-x))}
Hs = {}

for i in xrange(len(hidden_layer_sizes)):
    X = D if i == 0 else Hs[i-1]
    fan_in = X.shape[1]
    fan_out = hidden_layer_sizes[i]
    W = np.random.randn(fan_in, fan_out) * 0.01

    H = np.dot(X, W)
    H = act[nonlinearities[i]](H)
    Hs[i] = H

```

We will now consider a feedforward network with:

- input: 1000 points, each  $\in R^{500}$

```

D = np.random.randn(1000,500)
hidden_layer_sizes = [500]*10
nonlinearities = ['tanh']*len(hidden_layer_sizes)

act = {'relu':lambda x: np.maximum(0, x), 'tanh': lambda x: np.tanh(x),
       'sigmoid':lambda x: 1/(1 + np.exp(-x))}
Hs = {}


for i in xrange(len(hidden_layer_sizes)):
    X = D if i == 0 else Hs[i-1]
    fan_in = X.shape[1]
    fan_out = hidden_layer_sizes[i]
    W = np.random.randn(fan_in, fan_out) * 0.01

    H = np.dot(X, W)
    H = act[nonlinearities[i]](H)
    Hs[i] = H

```

We will now consider a feedforward network with:

- input: 1000 points, each  $\in R^{500}$
- input data is drawn from unit Gaussian



```

D = np.random.randn(1000,500)
hidden_layer_sizes = [500]*10
nonlinearities = ['tanh']*len(hidden_layer_sizes)

act = {'relu':lambda x: np.maximum(0, x), 'tanh': lambda x: np.tanh(x),
       'sigmoid':lambda x: 1/(1 + np.exp(-x))}
Hs = {}


for i in xrange(len(hidden_layer_sizes)):
    X = D if i == 0 else Hs[i-1]
    fan_in = X.shape[1]
    fan_out = hidden_layer_sizes[i]
    W = np.random.randn(fan_in, fan_out) * 0.01

    H = np.dot(X, W)
    H = act[nonlinearities[i]](H)
    Hs[i] = H

```

We will now consider a feedforward network with:

- input: 1000 points, each  $\in R^{500}$
- input data is drawn from unit Gaussian



- the network has 5 layers

```

D = np.random.randn(1000,500)
hidden_layer_sizes = [500]*10
nonlinearities = ['tanh']*len(hidden_layer_sizes)

act = {'relu':lambda x: np.maximum(0, x), 'tanh': lambda x: np.tanh(x),
       'sigmoid':lambda x: 1/(1 + np.exp(-x))}
Hs = {}


for i in xrange(len(hidden_layer_sizes)):
    X = D if i == 0 else Hs[i-1]
    fan_in = X.shape[1]
    fan_out = hidden_layer_sizes[i]
    W = np.random.randn(fan_in, fan_out) * 0.01

    H = np.dot(X, W)
    H = act[nonlinearities[i]](H)
    Hs[i] = H

```

We will now consider a feedforward network with:

- input: 1000 points, each  $\in R^{500}$
- input data is drawn from unit Gaussian



- the network has 5 layers
- each layer has 500 neurons



```

D = np.random.randn(1000,500)
hidden_layer_sizes = [500]*10
nonlinearities = ['tanh']*len(hidden_layer_sizes)

act = {'relu':lambda x: np.maximum(0, x), 'tanh': lambda x: np.tanh(x),
'sigmoid':lambda x: 1/(1 + np.exp(-x))}
Hs = {}


for i in xrange(len(hidden_layer_sizes)):
    X = D if i == 0 else Hs[i-1]
    fan_in = X.shape[1]
    fan_out = hidden_layer_sizes[i]
    W = np.random.randn(fan_in, fan_out) * 0.01

    H = np.dot(X, W)
    H = act[nonlinearities[i]](H)
    Hs[i] = H

```

We will now consider a feedforward network with:

- input: 1000 points, each  $\in R^{500}$
- input data is drawn from unit Gaussian




- the network has 5 layers
- each layer has 500 neurons
- we will run forward propagation on this network with different weight initializations

```
W = np.random.randn(fan_in, fan_out) * 0.01
```

- Let's try to initialize the weights to small random numbers

```
W = np.random.randn(fan_in, fan_out) * 0.01
```


- Let's try to initialize the weights to small random numbers
- We will see what happens to the activation across different layers



tanh activation functions

```
W = np.random.randn(fan_in, fan_out) * 0.01
```

- Let's try to initialize the weights to small random numbers
- We will see what happens to the activation across different layers




sigmoid activation functions

- What will happen during back propagation?


- What will happen during back propagation?
- Recall that  $\nabla w_1$  is proportional to the activation passing through it

- What will happen during back propagation?
- Recall that  $\nabla w_1$  is proportional to the activation passing through it
- If all the activations in a layer are very close to 0, what will happen to the gradient of the weights connecting this layer to the next layer?



- What will happen during back propagation?
- Recall that  $\nabla w_1$  is proportional to the activation passing through it
- If all the activations in a layer are very close to 0, what will happen to the gradient of the weights connecting this layer to the next layer?






- What will happen during back propagation?
- Recall that  $\nabla w_1$  is proportional to the activation passing through it
- If all the activations in a layer are very close to 0, what will happen to the gradient of the weights connecting this layer to the next layer?
- They will all be close to 0 (vanishing gradient problem)

```
W = np.random.randn(fan_in, fan_out)
```

- Let us try to initialize the weights to large random numbers

```
W = np.random.randn(fan_in, fan_out)
```


- Let us try to initialize the weights to large random numbers



tanh activation with large weights

```
W = np.random.randn(fan_in, fan_out)
```


- Let us try to initialize the weights to large random numbers



sigmoid activations with large weights

```
W = np.random.randn(fan_in, fan_out)
```


- Let us try to initialize the weights to large random numbers
- Most activations have saturated



sigmoid activations with large weights

```
W = np.random.randn(fan_in, fan_out)
```


- Let us try to initialize the weights to large random numbers
- Most activations have saturated
- What happens to the gradients at saturation?



sigmoid activations with large weights

```
W = np.random.randn(fan_in, fan_out)
```

- Let us try to initialize the weights to large random numbers
- Most activations have saturated
- What happens to the gradients at saturation?
- They will all be close to 0 (vanishing gradient problem)




sigmoid activations with large weights

- Let us try to arrive at a more principled way of initializing weights




- Let us try to arrive at a more principled way of initializing weights



$$s_{11} = \sum_{i=1}^n w_{1i} x_i$$

- Let us try to arrive at a more principled way of initializing weights



•


•

 $s_{1n}$ 

$$s_{11} = \sum_{i=1}^n w_{1i} x_i$$

$$Var(s_{11}) = Var\left(\sum_{i=1}^n w_{1i} x_i\right) = \sum_{i=1}^n Var(w_{1i} x_i)$$

- Let us try to arrive at a more principled way of initializing weights




$$s_{11} = \sum_{i=1}^n w_{1i} x_i$$

$$Var(s_{11}) = Var\left(\sum_{i=1}^n w_{1i} x_i\right) = \sum_{i=1}^n Var(w_{1i} x_i)$$

$$= \sum_{i=1}^n \left[ (E[w_{1i}])^2 Var(x_i) + (E[x_i])^2 Var(w_{1i}) + Var(x_i) Var(w_{1i}) \right]$$

- Let us try to arrive at a more principled way of initializing weights



$$s_{11} = \sum_{i=1}^n w_{1i}x_i$$


$$Var(s_{11}) = Var\left(\sum_{i=1}^n w_{1i}x_i\right) = \sum_{i=1}^n Var(w_{1i}x_i)$$

$$= \sum_{i=1}^n [(E[w_{1i}])^2 Var(x_i)$$

$$+ (E[x_i])^2 Var(w_{1i}) + Var(x_i)Var(w_{1i})]$$

- [Assuming 0 Mean inputs and weights]

- Let us try to arrive at a more principled way of initializing weights



$$s_{11} = \sum_{i=1}^n w_{1i}x_i$$


$$Var(s_{11}) = Var\left(\sum_{i=1}^n w_{1i}x_i\right) = \sum_{i=1}^n Var(w_{1i}x_i)$$

$$= \sum_{i=1}^n [(E[w_{1i}])^2 Var(x_i)$$

$$+ (E[x_i])^2 Var(w_{1i}) + Var(x_i)Var(w_{1i})]$$

- [Assuming 0 Mean inputs and weights]
- [Assuming  $Var(x_i) = Var(x) \forall i$ ]

- Let us try to arrive at a more principled way of initializing weights



$$s_{11} = \sum_{i=1}^n w_{1i} x_i$$


$$Var(s_{11}) = Var\left(\sum_{i=1}^n w_{1i} x_i\right) = \sum_{i=1}^n Var(w_{1i} x_i)$$

$$= \sum_{i=1}^n \left[ (E[w_{1i}])^2 Var(x_i) + (E[x_i])^2 Var(w_{1i}) + Var(x_i) Var(w_{1i}) \right]$$

$$= \sum_{i=1}^n Var(x_i) Var(w_{1i})$$

- [Assuming 0 Mean inputs and weights]
- [Assuming  $Var(x_i) = Var(x) \forall i$ ]

- Let us try to arrive at a more principled way of initializing weights




$$s_{11} = \sum_{i=1}^n w_{1i} x_i$$

$$\begin{aligned} Var(s_{11}) &= Var\left(\sum_{i=1}^n w_{1i} x_i\right) = \sum_{i=1}^n Var(w_{1i} x_i) \\ &= \sum_{i=1}^n \left[ (E[w_{1i}])^2 Var(x_i) \right. \\ &\quad \left. + (E[x_i])^2 Var(w_{1i}) + Var(x_i) Var(w_{1i}) \right] \\ &= \sum_{i=1}^n Var(x_i) Var(w_{1i}) \end{aligned}$$

- [Assuming 0 Mean inputs and weights]
- [Assuming  $Var(x_i) = Var(x) \forall i$ ]
- [Assuming  $Var(w_{1i}) = Var(w) \forall i$ ]

- Let us try to arrive at a more principled way of initializing weights



$$s_{11} = \sum_{i=1}^n w_{1i} x_i$$


$$Var(s_{11}) = Var\left(\sum_{i=1}^n w_{1i} x_i\right) = \sum_{i=1}^n Var(w_{1i} x_i)$$

$$\begin{aligned} &= \sum_{i=1}^n \left[ (E[w_{1i}])^2 Var(x_i) \right. \\ &\quad \left. + (E[x_i])^2 Var(w_{1i}) + Var(x_i) Var(w_{1i}) \right] \\ &= \sum_{i=1}^n Var(x_i) Var(w_{1i}) \\ &= (n Var(w))(Var(x)) \end{aligned}$$


- [Assuming 0 Mean inputs and weights]
- [Assuming  $Var(x_i) = Var(x) \forall i$ ]
- [Assuming  $Var(w_{1i}) = Var(w) \forall i$ ]



- In general,




$$Var(S_{1i}) = (nVar(w))(Var(x))$$



- In general,

$$Var(S_{1i}) = (nVar(w))(Var(x))$$


- What would happen if  $nVar(w) \gg 1$  ?



- In general,

$$Var(S_{1i}) = (nVar(w))(Var(x))$$


- What would happen if  $nVar(w) \gg 1$ ?
- The variance of  $S_{1i}$  will be large



- In general,

$$Var(S_{1i}) = (nVar(w))(Var(x))$$

- What would happen if  $nVar(w) \gg 1$ ?
- The variance of  $S_{1i}$  will be large
- What would happen if  $nVar(w) \rightarrow 0$ ?




- In general,


$$Var(S_{1i}) = (nVar(w))(Var(x))$$

- What would happen if  $nVar(w) \gg 1$  ?
- The variance of  $S_{1i}$  will be large
- What would happen if  $nVar(w) \rightarrow 0$ ?
- The variance of  $S_{1i}$  will be small

- Let us see what happens if we add one more layer




- Let us see what happens if we add one more layer
- Using the same procedure as above we will arrive at



- Let us see what happens if we add one more layer
- Using the same procedure as above we will arrive at


$$Var(s_{21}) = \sum_{i=1}^n Var(s_{1i})Var(w_{2i})$$





- Let us see what happens if we add one more layer
- Using the same procedure as above we will arrive at

$$\begin{aligned} Var(s_{21}) &= \sum_{i=1}^n Var(s_{1i}) Var(w_{2i}) \\ &= n Var(s_{1i}) Var(w_2) \end{aligned}$$




$$\boxed{Var(S_{i1}) = nVar(w_1)Var(x)}$$

- Let us see what happens if we add one more layer
- Using the same procedure as above we will arrive at

$$Var(s_{21}) = \sum_{i=1}^n Var(s_{1i})Var(w_{2i})$$

$$= nVar(s_{1i})Var(w_2)$$




$$\boxed{Var(S_{i1}) = nVar(w_1)Var(x)}$$

- Let us see what happens if we add one more layer
- Using the same procedure as above we will arrive at

$$Var(s_{21}) = \sum_{i=1}^n Var(s_{1i})Var(w_{2i})$$

$$= nVar(s_{1i})Var(w_2)$$

$$Var(s_{21}) \propto [nVar(w_2)][nVar(w_1)]Var(x)$$




$$\boxed{Var(S_{i1}) = nVar(w_1)Var(x)}$$

- Let us see what happens if we add one more layer
- Using the same procedure as above we will arrive at

$$\begin{aligned} Var(s_{21}) &= \sum_{i=1}^n Var(s_{1i})Var(w_{2i}) \\ &= nVar(s_{1i})Var(w_2) \end{aligned}$$

$$Var(s_{21}) \propto [nVar(w_2)][nVar(w_1)]Var(x)$$

$$\propto [nVar(w)]^2Var(x)$$



$$\boxed{Var(S_{i1}) = nVar(w_1)Var(x)}$$

- Let us see what happens if we add one more layer
- Using the same procedure as above we will arrive at

$$\begin{aligned} Var(s_{21}) &= \sum_{i=1}^n Var(s_{1i})Var(w_{2i}) \\ &= nVar(s_{1i})Var(w_2) \end{aligned}$$

$$Var(s_{21}) \propto [nVar(w_2)][nVar(w_1)]Var(x)$$

$$\propto [nVar(w)]^2Var(x)$$

Assuming weights across all layers  
have the same variance

- In general,

- In general,

$$\text{Var}(s_{ki}) = [n\text{Var}(w)]^k \text{Var}(x)$$

- In general,

$$\text{Var}(s_{ki}) = [n\text{Var}(w)]^k \text{Var}(x)$$

- To ensure that variance in the output of any layer does not blow up or shrink we want:

$$n \text{Var}(w) = 1$$



- In general,

$$\text{Var}(s_{ki}) = [n\text{Var}(w)]^k \text{Var}(x)$$

- To ensure that variance in the output of any layer does not blow up or shrink we want:

$$n \text{Var}(w) = 1$$

- If we draw the weights from a unit Gaussian and scale them by  $\frac{1}{\sqrt{n}}$  then, we have :

- In general,

$$\text{Var}(s_{ki}) = [n\text{Var}(w)]^k \text{Var}(x)$$

- To ensure that variance in the output of any layer does not blow up or shrink we want:

$$n \text{Var}(w) = 1$$

- If we draw the weights from a unit Gaussian and scale them by  $\frac{1}{\sqrt{n}}$  then, we have :

$$= n \text{Var}\left(\frac{w}{\sqrt{n}}\right)$$

$$\boxed{\text{Var}(aw) = a^2(\text{Var}(w))}$$

- In general,

$$\text{Var}(s_{ki}) = [n\text{Var}(w)]^k \text{Var}(x)$$

- To ensure that variance in the output of any layer does not blow up or shrink we want:

$$n \text{Var}(w) = 1$$

- If we draw the weights from a unit Gaussian and scale them by  $\frac{1}{\sqrt{n}}$  then, we have :

$$= n \text{Var}\left(\frac{w}{\sqrt{n}}\right)$$

$$= n * \frac{1}{n} \text{Var}(w) = 1$$

$$\boxed{Var(aw) = a^2(Var(w))}$$

- In general,

$$Var(s_{ki}) = [nVar(w)]^k Var(x)$$

- To ensure that variance in the output of any layer does not blow up or shrink we want:

$$n Var(w) = 1$$

- If we draw the weights from a unit Gaussian and scale them by  $\frac{1}{\sqrt{n}}$  then, we have :

$$= n Var\left(\frac{w}{\sqrt{n}}\right)$$


$$= n * \frac{1}{n} Var(w) = 1 \leftarrow (\text{Unit Gaussian})$$

```
W = np.random.randn(fan_in, fan_out) / sqrt(fan_in)
```

- Let's see what happens if we use this initialization

```
W = np.random.randn(fan_in, fan_out) / sqrt(fan_in)
```


- Let's see what happens if we use this initialization



tanh activation


```
W = np.random.randn(fan_in, fan_out) / sqrt(fan_in)
```

- Let's see what happens if we use this initialization




sigmoid activations


- However this does not work for ReLU neurons






- However this does not work for ReLU neurons
- Why ?





- However this does not work for ReLU neurons
- Why ?
- Intuition: *He et.al.* argue that a factor of 2 is needed when dealing with ReLU Neurons



- However this does not work for ReLU neurons
- Why ?
- Intuition: *He et.al.* argue that a factor of 2 is needed when dealing with ReLU Neurons
- Intuitively this happens because the range of ReLU neurons is restricted only to the positive half of the space

```
W = np.random.randn(fan_in, fan_out) / sqrt(fan_in/2)
```

- Indeed when we account for this factor of 2 we see better performance

```
W = np.random.randn(fan_in, fan_out) / sqrt(fan_in/2)
```

- Indeed when we account for this factor of 2 we see better performance

